



# Loan Default Predictor

<https://github.com/yxing12/loan-default-predictor>

Yunfei (Cynthia) Xing  
Data Science Institute

December 6, 2023

A low-angle, upward-looking photograph of several modern skyscrapers with glass and steel facades. The buildings are set against a clear blue sky with a few wispy clouds. A thin white rectangular frame is superimposed over the center of the image, enclosing the text.

01

Project  
Recap

# Recap



## Problem

Loan Default - situation where a borrower fails to repay a loan according to the terms agreed upon.

Negatively impacts the borrower (credit rating) and the financial institution (revenue, reputation).

**Leverage ML techniques to predict loan defaults.**

## Dataset

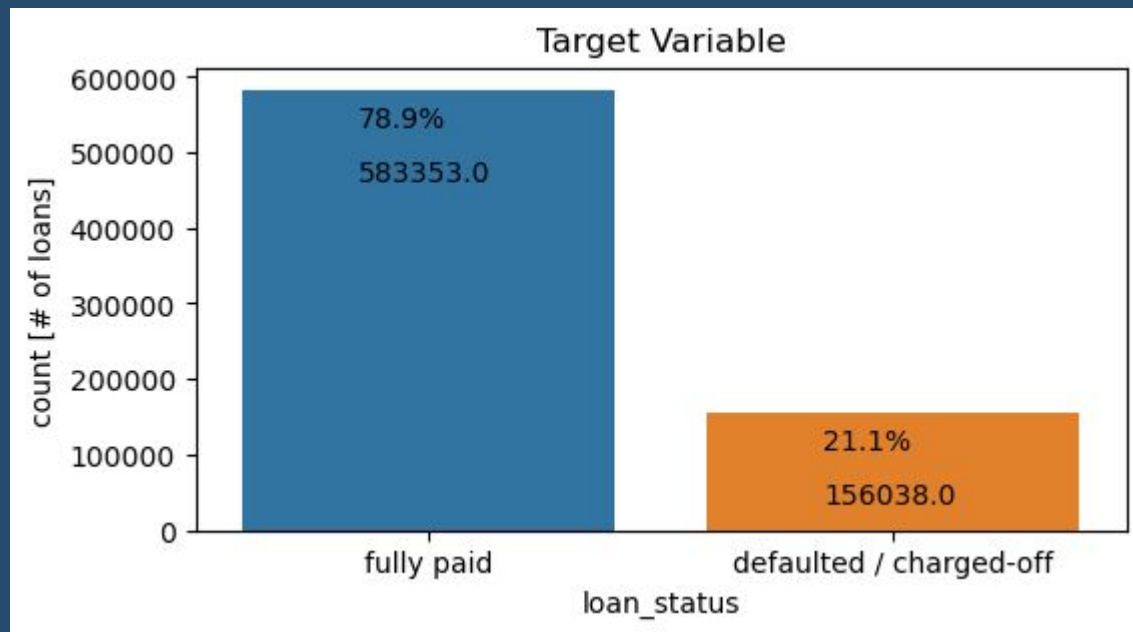


- Lending Club's loan data from 2007-2017 adopted from [Kaggle](#)
  - Aggregated data from [Lending Club website](#)
- Cleaned Dataset shape: (739,391, 32)
- I.I.D: Yes
- **Classification Problem** - categorize the loan into one of the two class (fully paid/defaulted) based on loan specs and borrower financials

## Recap - Target Variable

- ~80% loans fully paid  
~20% charged off /  
defaulted

→ **unbalanced  
classification problem**



\* A charge-off is a debt that a creditor has given up trying to collect on after borrower have missed payments for several months.

# Recap - Splitting

Basic train\_test\_split

- Large i.i.d dataset (>700k rows)
- **98-1-1 train-val-test split**

```
train + val shape shape: (731997, 31) (731997,)
```

```
test shape: (7394, 31) (7394,)
```

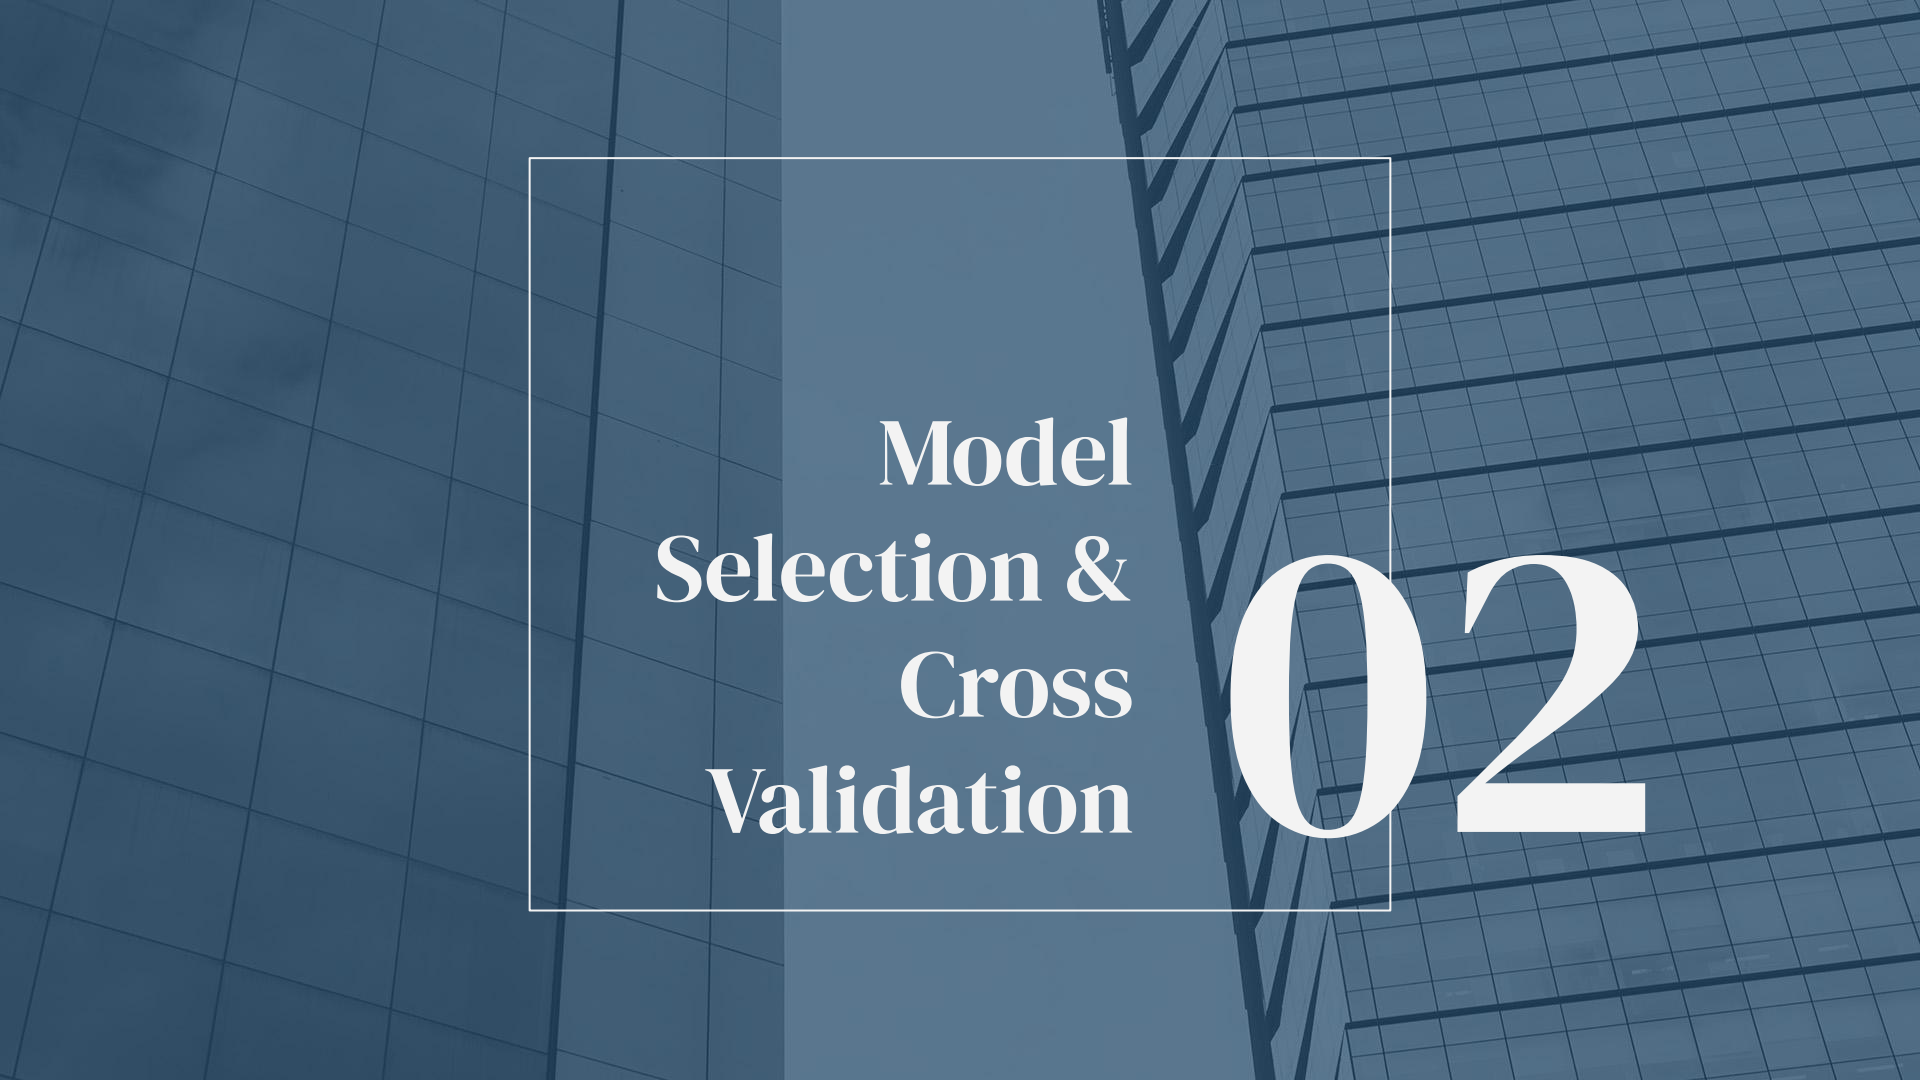
```
y_other Distribution: Fully-Paid: 78.90%, Defaulted/Charged-off: 21.10%
```

```
y_test Distribution: Fully-Paid: 78.90%, Defaulted/Charged-off: 21.10%
```

Preprocessing

- OneHotEncoder
  - categorical
- OrdinalEncoder
  - ordinal
- MinMaxScaler
- StandardScaler
  - continuous

```
# collect the features
cat_ftrs = ['home_ownership','verification_status','purpose',\
            'addr_state','initial_list_status','application_type']
ordinal_ftrs = ['grade','sub_grade']
ordinal_cats = [['A','B','C','D','E','F','G'],\
                ['A1','A2','A3','A4','A5',\
                 'B1','B2','B3','B4','B5',\
                 'C1','C2','C3','C4','C5',\
                 'D1','D2','D3','D4','D5',\
                 'E1','E2','E3','E4','E5','F1','F2','F3','F4','F5',\
                 'G1','G2','G3','G4','G5']]
max_min_ftrs = ['loan_amnt','int_rate','installment','fico_score']
std_num_ftrs = ['term','dti','earliest_cr_line','open_acc','pub_rec',\
                'revol_util','total_pymnt','total_rec_int','last_pymnt_amnt',\
                'acc_open_past_24mths','avg_cur_bal','bc_open_to_buy','bc_util',\
                'mo_sin_old_rev_tl_op','mo_sin_rcnt_rev_tl_op','mort_acc',\
                'num_actv_rev_tl','pub_rec_bankruptcies','log_annual_inc']
```



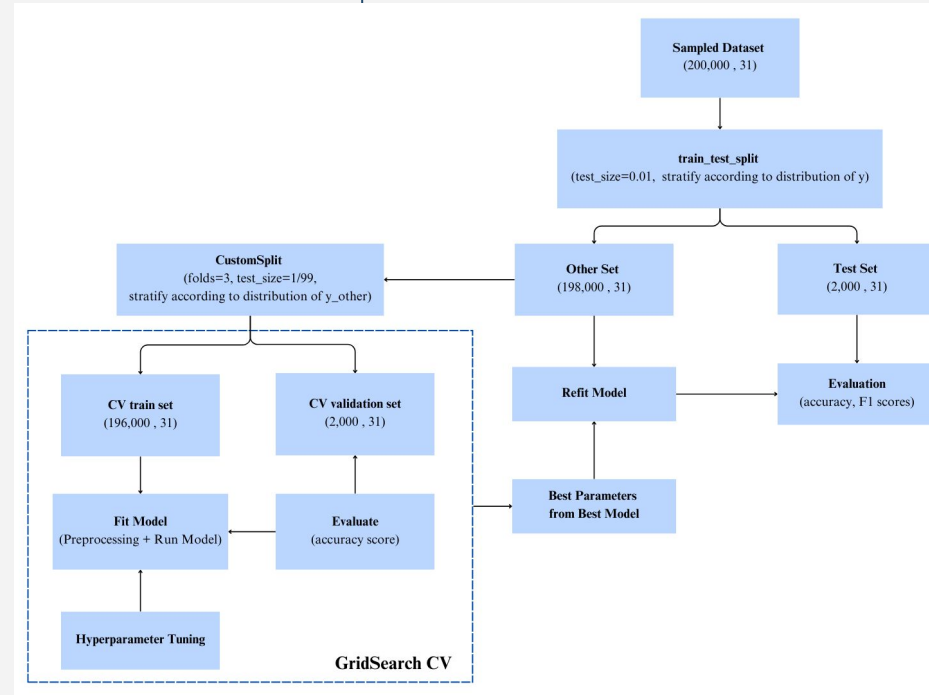
# Model Selection & Cross Validation

# 02



# Cross Validation Pipeline

- Define a “CustomSplit” class
  - Functions like StratifiedKFold
  - Configurable number of splits for CV
  - Allows setting a specific proportion for test data (98-1-1)
- CV Pipeline
  - Sampled 200,000 data
  - Iterate through 5 random states
  - Split the data into 'other' and 'test'
  - Create pipeline and GridSearchCV
    - preprocess the data
    - undergo GridSearchCV to train and validate the model
    - scoring = accuracy
  - Find best model, refit on full dataset, and calculate test scores



# ML Algos & Hyperparameters

**'logisticregression\_\_C':**  
[0.0001, 0.001, 0.01, 0.1, 1]

**Logistic  
Regression**

**'xgbclassifier\_\_reg\_alpha':**  
[1e-1, 1e0, 1e1],

**'xgbclassifier\_\_reg\_lambda':**  
[1e-1, 1e0, 1e1],

**'xgbclassifier\_\_max\_depth':**  
[5, 10, 20]

\* Redefined Pipeline to  
implement early\_stopping,  
used  
early\_stopping\_rounds=10

**XGB Classifier**

**'randomforestclassifier\_\_  
max\_depth':** [1, 3, 10, 30,  
100],

**'randomforestclassifier\_\_  
max\_features':** [0.25, 0.5,  
0.75, 1.0]

**Random Forest  
Classifier**

**'kneighborsclassifier\_\_n\_  
neighbors':** [10, 20, 30,  
100],

**'kneighborsclassifier\_\_w  
eights':** ['uniform',  
'distance']

**KNeighbors  
Classifier**

**Tuned on a subsample of 200,000 data points for 5 random states**



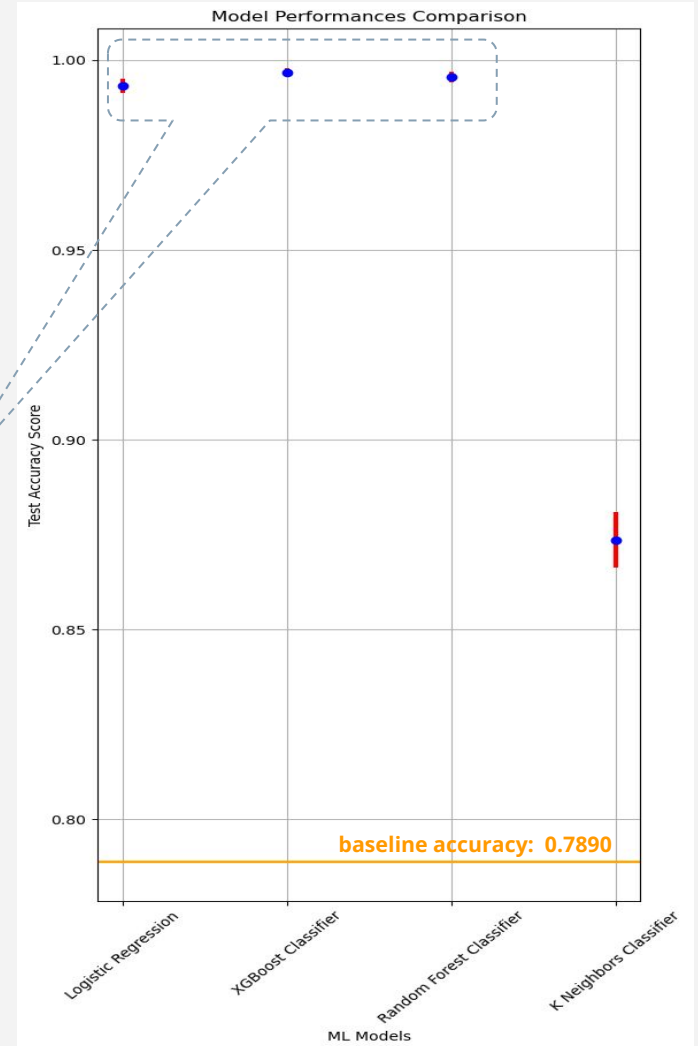
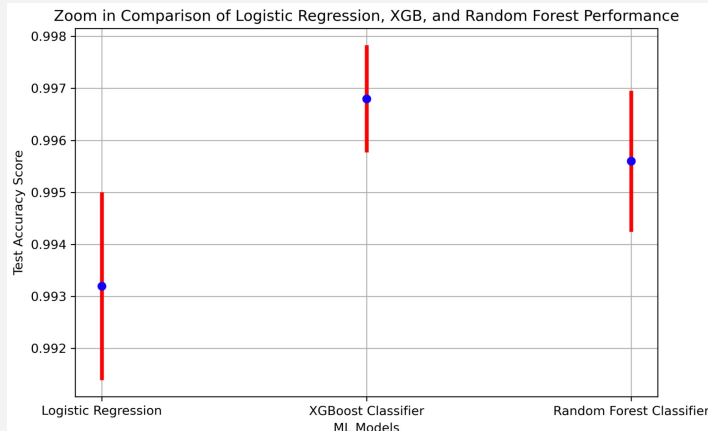
A low-angle, upward-looking photograph of several modern skyscrapers with glass facades, set against a blue sky with scattered white clouds. The perspective makes the buildings appear to converge towards the top of the frame. A white rectangular border is superimposed over the center of the image, framing the text.

03

# Results & Outlook

# Results from CV

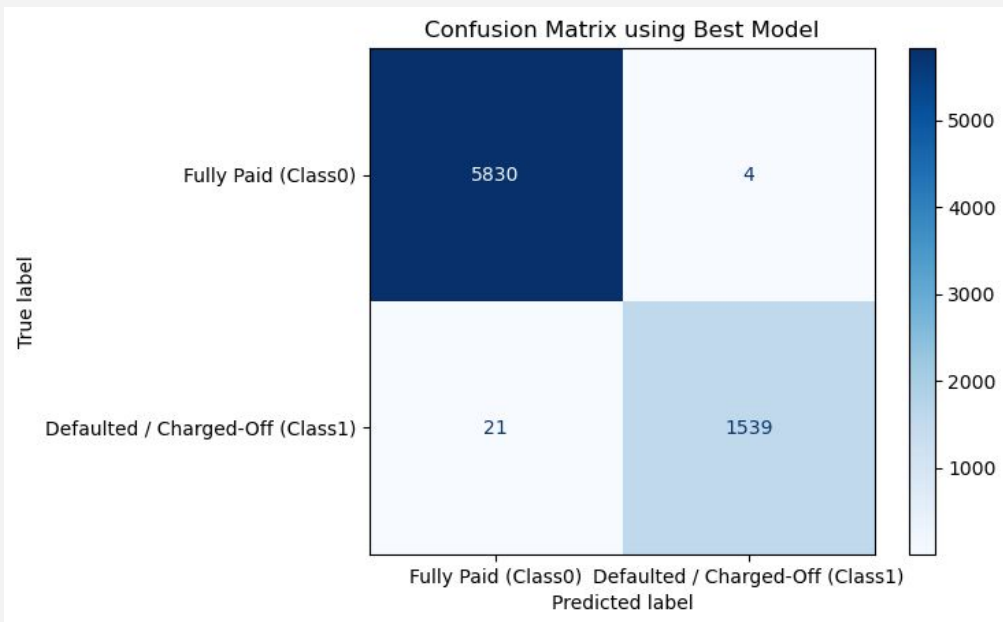
	Model	Mean Test Score	Std Test Score	#stds Above Baseline Accuracy
0	Logistic Regression	0.9932	0.001806	113.0
1	XGBoost Classifier	0.9968	0.001030	202.0
2	Random Forest Classifier	0.9956	0.001356	152.0
3	K Neighbors Classifier	0.8736	0.007317	12.0



- KNN struggles in high-dimensional spaces (the curse of dimensionality) and with non-linear relationships between features
- KNN can be significantly impacted by noise and imbalances in the dataset

# Best Model

**XGBClassifier (random\_state=42, colsample\_bytree=0.9, subsample=0.8, max\_depth=20, reg\_alpha=0.1, reg\_lambda=1.0, early\_stopping\_rounds=10)**

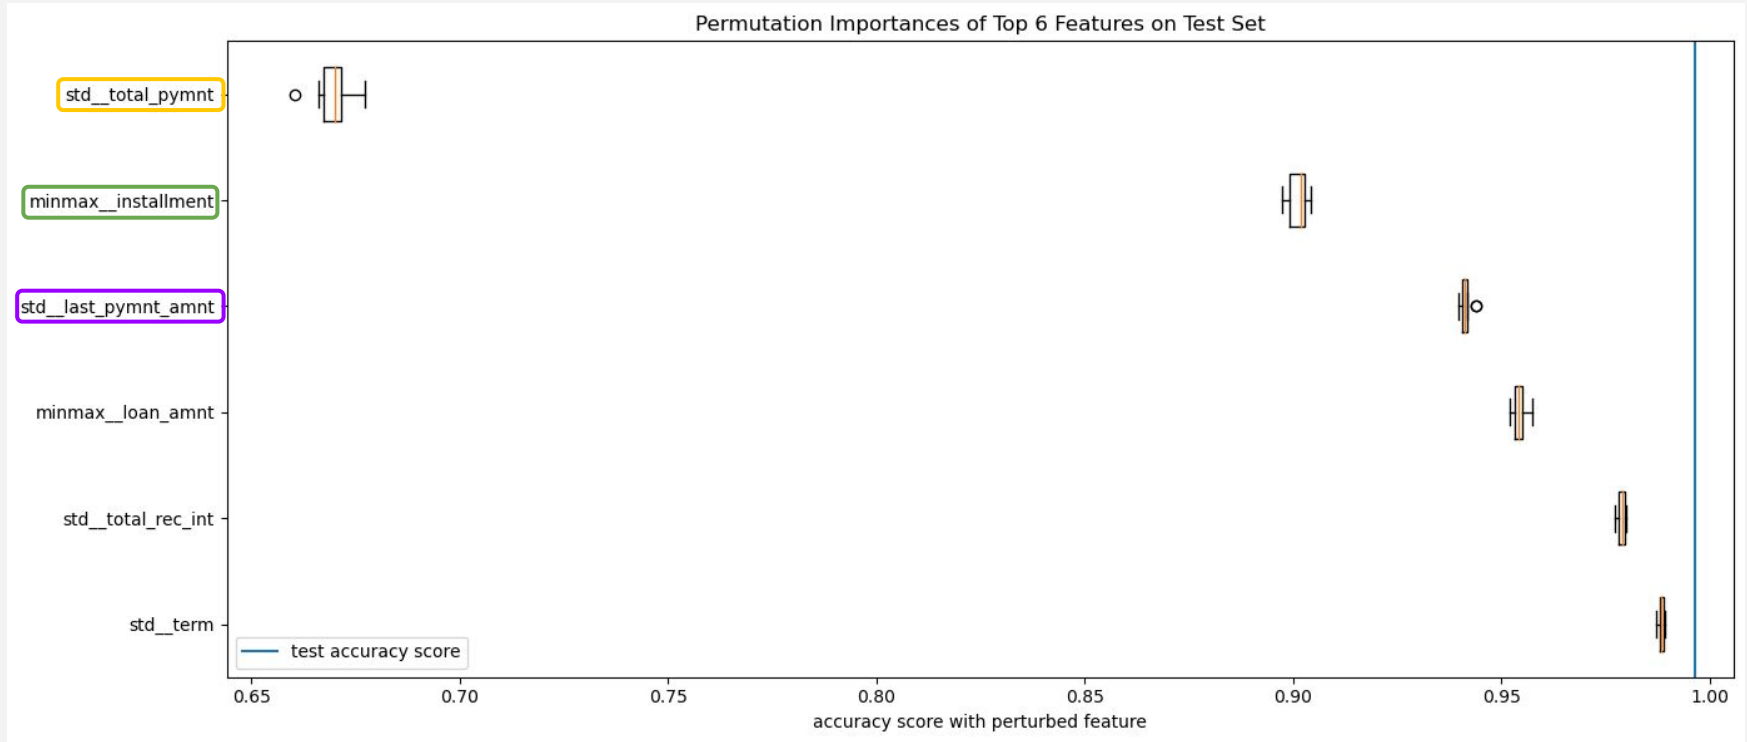


- Refit best model on full dataset
- 98-1-1 split ratio
- cm threshold = 0.2

```
Validation accuracy: 0.9981065728969435
Validation precision: 0.9980670103092784
Validation recall: 0.992948717948718
Validation F1 score: 0.9955012853470437

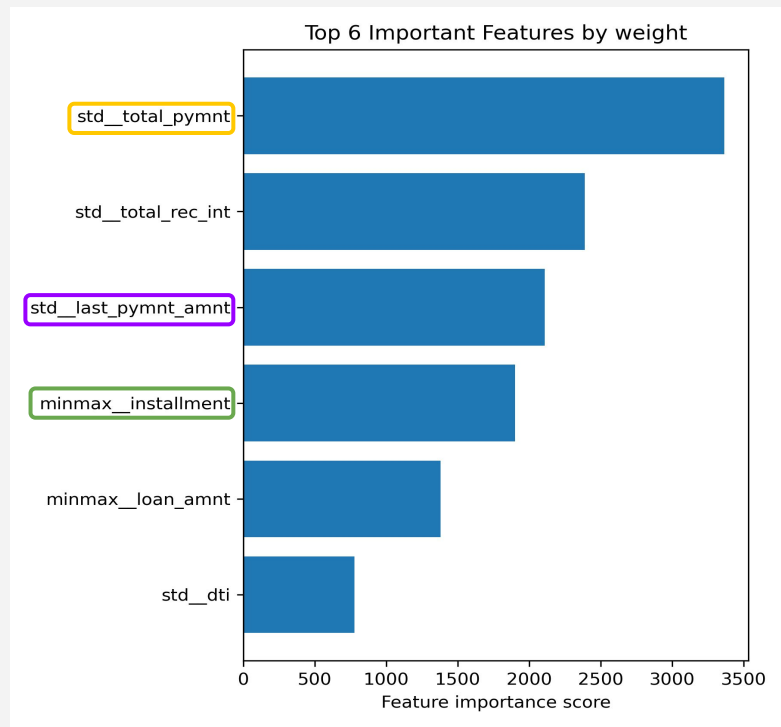
Test accuracy: 0.9966188801731133
Test precision: 0.9974076474400518
Test recall: 0.9865384615384616
Test F1 score: 0.9919432806961006
```

# Global Feature Importance

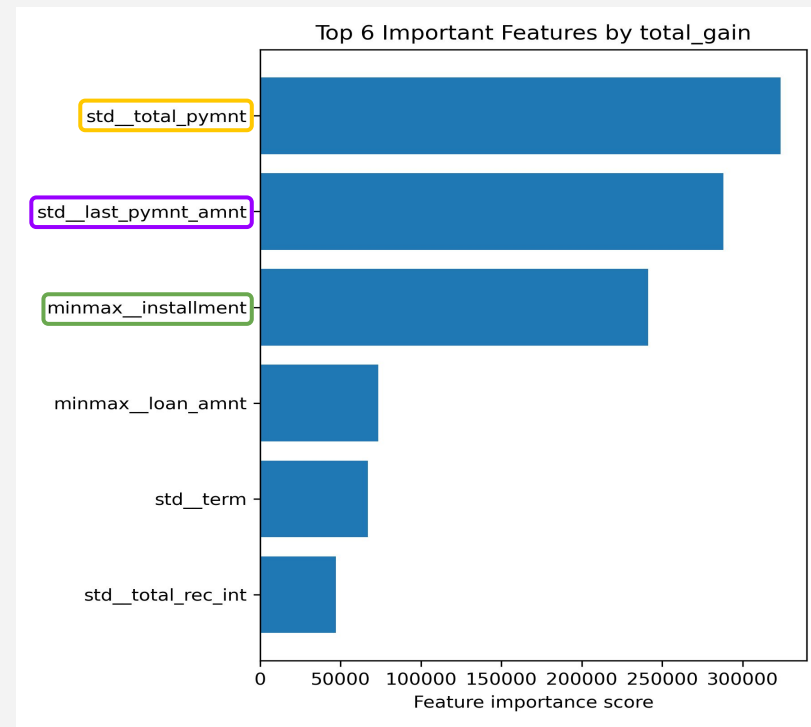


estimates predictive influence based on the drop in model performance due to random perturbation of a feature

# Global Feature Importance




Weight - number of times a feature is used to split the data across all trees.



Total Gain - relative contribution of the corresponding feature to the model's accuracy

# Global Feature Importance



estimates predictive influence based on Shapley values from game theory

# Local Feature Importance


point index = 234

higher  $\leftrightarrow$  lower

$f(x)$

base value

-5.82




point index = 520

base value

higher  $\leftrightarrow$  lower

$f(x)$

0.27




point index = 777

higher  $\leftrightarrow$  lower

$f(x)$

base value

-7.81



gauge the contribution of a feature to the predictability of a certain point



# Outlook

Given more time and more computing power, the project could be improved by:

## Increasing CV folds



Due to limitation in computing power, I only used 3 folds in CV. Increasing to 5 or 10 folds would be ideal.

## Comprehensive Dataset & Hyperparameter Tuning



Using the full dataset for CV, along with tuning a more thorough hyperparameters could significantly refine the model's generalizability.

## Try out more models and techniques



Due to the large size of the dataset, I was unable to run SVM because it was taking too long. Also, I'm curious about exploring non-ML techniques, such as Neural Networks, to see how they can influence model performance and interpretability.

## Misclassification Reduction



I would like to explore more ways to minimize false negatives and false positives, which is especially important in loan default prediction. For instance, false negatives can cause credit issues for both borrowers and financial institutions.



**Thank you!**

**Questions?**